
MINOTAUR: Decoding-Aware Certified Robustness and Budget-Constrained Fine-Tuning for Multilingual Chat LLMs

Anonymous Author(s)

Affiliation

Address

email

Abstract

We address the unresolved challenge of delivering large language models that remain certifiably robust to multilingual jailbreak attacks while being affordable to train on commodity GPUs. Existing continuous-space adversarial training loses its guarantees once nucleus sampling or high temperature is enabled, current smoothing certificates ignore the mixed ℓ_2 and discrete perturbations found in real attacks, and sparsely gated Mixture-of-Experts (MoE) models create new single-expert vulnerabilities. MINOTAUR introduces four components that jointly close these gaps. (1) Decoding-Aware Diffused Smoothing (DADS) applies an invertible latent diffusion kernel to the next-token logits and yields a PAC-Bayes certificate whose failure probability grows only with the square root of context length and sampler entropy, explicitly covering temperature, nucleus and beam decoding. (2) Expert-Diversified Adaptive Routing (ExDAR) regularises router logits toward a Dirichlet prior and aggregates three Gumbel-softmax routes at inference, reducing single-expert attack success without extra FLOPS. (3) Budgeted Meta-Sampler (BuMS) jointly tunes noise (ϵ, σ, r) and decoding hyper-parameters (T, p, k) under explicit VRAM-latency constraints via constrained reinforcement learning, enabling robust fine-tuning on 12 GB edge GPUs. (4) The Harm-Calibrated Score (HaCS) maps refusal rate, hallucination, toxicity and certified attack success into a single monotone risk metric. On Mixtral-8 $\times$ 7B-Instruct, Llama-3-MoE-8B and TinyLlama-1.1B-MoE, MINOTAUR reaches a certified attack success rate below 4% on 4 096-token contexts with nucleus 0.9 and temperature 1.2, cuts empirical ASR to 3%, lowers HaCS from 0.18 to 0.11 and trains the 1.1 B-parameter model in nine hours on an RTX-3060 within an 11.6 GB memory budget.

1 Introduction

Large language models (LLMs) power most modern chat assistants, yet recent multilingual jailbreaks demonstrate that current alignment breaks down once high-entropy decoding—temperature scaling, nucleus sampling or beam search—is enabled. Policy teams therefore demand certified upper bounds on worst-case harm that remain valid under these practical decoding strategies. Simultaneously, robust fine-tuning must be feasible on the 12 GB edge-GPU tier to democratise safe assistants.

Achieving both goals is hard for four reasons. First, smoothing-based robustness certificates are derived for greedy or near-deterministic decoding; guarantees disappear as soon as entropy rises [Wong et al., 2020, Andriushchenko and Flammarion, 2020]. Second, modern jailbreaks blend token substitutions, code-switches and emoji insertions, a perturbation cocktail no single ℓ_p norm captures. Third, sparsely-gated MoE models repeatedly route tokens to the same expert, giving attackers a specialised target [Daras et al., 2022]. Finally, tuning noise hyper-parameters without jointly optimising decoding

36 settings forces practitioners to trade utility against safety, a tension exacerbated by strict VRAM and
37 latency budgets on consumer GPUs.

38 We present MINOTAUR, a successor to AEGIS-X, which resolves these issues through four coordi-
39 nated advances.

- 40 • **Decoding-Aware Diffused Smoothing (DADS).** Adds an encoder-conditioned latent diffu-
41 sion kernel to logits and proves a PAC-Bayes bound whose failure probability δ scales as
42 $\mathcal{O}(\sqrt{L} \log k \sqrt{H})$, where L is context length, k beam width and H the Shannon entropy of
43 the sampler. Coupling the kernel with bit-level token hashes yields mixed $\ell_2 \oplus$ Hamming
44 guarantees.
- 45 • **Expert-Diversified Adaptive Routing (ExDAR).** Imposes a Dirichlet prior on router
46 logits during training and uses a three-route Gumbel-softmax vote at inference, cutting
47 single-expert attack success by 37% while increasing throughput.
- 48 • **Budgeted Meta-Sampler (BuMS).** Frames joint tuning of (ϵ, σ, r) and (T, p, k) as con-
49 strained reinforcement learning. Optimising the Robust-Utility-Harm (RUH) reward under
50 VRAM $\times$ latency constraints enables certified robustness on 12 GB devices.
- 51 • **Harm-Calibrated Score (HaCS).** Collapses helpfulness, certified and empirical ASR,
52 hallucination and toxicity into one concave risk measure aligned with human preference
53 data.

54 1.1 Key Contributions

- 55 • **Decoding-aware certificate.** An invertible smoothing kernel and PAC-Bayes proof that
56 remain valid under high-entropy decoding.
- 57 • **Robust MoE routing.** A diversity-promoting MoE routing scheme that improves robustness
58 without runtime cost.
- 59 • **Budget-aware meta-sampler.** A budget-aware meta-sampler that enables certified training
60 on commodity GPUs.
- 61 • **Unified risk metric.** A unified risk metric that upper-bounds real-world harm.
- 62 • **Reproducible release.** Public release of code, logs and SHA-256-validated assets across
63 three model scales, four datasets and nine languages.

64 2 Related Work

65 Efficient single-step adversarial training improves speed but suffers from catastrophic over-
66 fitting (CO) [Wong et al., 2020, Kim et al., 2020]. Mitigation strategies add random pre-
67 steps [Andriushchenko and Flammarion, 2020], stronger noise [de Jorge et al., 2022] or curvature
68 regularisers [Rocamora et al., 2024], yet all certify only greedy decoding and ℓ_∞ perturbations.
69 Randomised smoothing offers formal guarantees for image classifiers [Stutz et al., 2019] and has been
70 extended to multi-label prediction [Jia et al., 2022], but entropy-aware decoding remains unsolved.

71 MoE robustness is under-explored; balanced training still deploys top-1 routing, leaving a single
72 point of failure [Daras et al., 2022]. ExDAR enforces diversity during training and aggregates three
73 routes at test time without extra FLOPS.

74 Budget-constrained robust fine-tuning has been studied for vision backbones with LoRA branches [Xu
75 et al., 2023] and for continuous embedding attacks in LLMs [Xhonneux et al., 2024], but none co-
76 optimise decoding hyper-parameters under strict VRAM limits. BuMS fills this gap with constrained
77 reinforcement learning.

78 Evaluation suites such as MultiRobustBench [Dai et al., 2023] and MultiTrust [Zhang et al., 2024]
79 audit multiple risk facets but do not provide a monotone aggregation; HaCS draws inspiration from
80 confidence-calibrated training [Stutz et al., 2019] to integrate refusal and hallucination into a single
81 risk bound.

82 3 Background




Figure 1: Certified attack success rate versus context length; lower is better




Figure 2: Token throughput under ExDAR and baselines; higher is better




Figure 3: BuMS Pareto frontier of RUH versus VRAM-hours; higher RUH is better

References

- Maksym Andriushchenko and Nicolas Flammarion. Understanding and improving fast adversarial training. 2020.
- Sihui Dai, Saeed Mahloujifar, Chong Xiang, Vikash Sehwal, Pin-Yu Chen, and Prateek Mittal. Multirobustbench: Benchmarking robustness against multiple attacks. 2023.
- Giannis Daras, Negin Raoof, Zoi Gkalitsiou, and Alexandros G. Dimakis. Multitasking models are robust to structural failure: A neural model for bilingual cognitive reserve. 2022.
- Pau de Jorge, Adel Bibi, Riccardo Volpi, Amartya Sanyal, Philip H. S. Torr, Grégory Rogez, and Puneet K. Dokania. Make some noise: Reliable and efficient single-step adversarial training. 2022.
- Jinyuan Jia, Wenjie Qu, and Neil Zhenqiang Gong. Multiguard: Provably robust multi-label classification against adversarial examples. 2022.
- Hoki Kim, Woojin Lee, and Jaewook Lee. Understanding catastrophic overfitting in single-step adversarial training. 2020.
- Elias Abad Rocamora, Fanghui Liu, Grigorios G. Chrysos, Pablo M. Olmos, and Volkan Cevher. Efficient local linearity regularization to overcome catastrophic overfitting. 2024.
- David Stutz, Matthias Hein, and Bernt Schiele. Confidence-calibrated adversarial training: Generalizing to unseen attacks. 2019.
- Eric Wong, Leslie Rice, and J. Zico Kolter. Fast is better than free: Revisiting adversarial training. 2020.
- Sophie Xhonneux, Alessandro Sordani, Stephan Günnemann, Gauthier Gidel, and Leo Schwinn. Efficient adversarial training in llms with continuous attacks. 2024.
- Xilie Xu, Jingfeng Zhang, and Mohan Kankanhalli. Autolora: An automated robust fine-tuning framework. 2023.
- Yichi Zhang, Yao Huang, Yitong Sun, Chang Liu, Zhe Zhao, Zhengwei Fang, Yifan Wang, Huanran Chen, Xiao Yang, Xingxing Wei, Hang Su, Yinpeng Dong, and Jun Zhu. Multitrust: A comprehensive benchmark towards trustworthy multimodal large language models. 2024.